

ABHISHEK MANOJ SUTARIA

✉ abhishek.sutaria@gmail.com | ☎ +1 (930) 333-7194 | [🌐 LinkedIn](#) | [🐙 GitHub](#) | [📁 Portfolio](#)

EDUCATION

Indiana University Bloomington
Master of Science in Data Science

Aug 2024 - May 2026
GPA 3.72/4.0

TECHNICAL SKILLS

Languages: Python, SQL, R, Spark (PySpark), Shell Scripting, Java, C++
ML & AI: Scikit-learn, XGBoost, TensorFlow, PyTorch, Keras, OpenCV, SHAP, Optuna, MLflow
GenAI: LangChain, LangGraph, CrewAI, RAG, Hugging Face, LlamaIndex, Sentence Transformers, Llama-2
Vector & Search: Pinecone, FAISS, Weaviate, Qdrant
Data Engineering: Airflow, Spark, Kafka, Hadoop, Polars, AWS Glue, AWS Lambda, AWS Kinesis
Cloud & DevOps: AWS (S3, EC2, EKS), GCP, Docker, Kubernetes, Git, GitHub, CI/CD, BitBucket
Databases: PostgreSQL, MySQL, MongoDB, Neo4j, Snowflake, BigQuery, DynamoDB, ClickHouse
APIs & Deployment: FastAPI, Streamlit, Gradio, WebRTC, REST APIs

WORK EXPERIENCE

Kelley School of Business, Indiana | *AI Research Engineer* Sep 2025 – Present

- Engineered a multi-agent RAG tutoring assistant using LangChain and CrewAI over multimodal academic data (PDFs, lecture transcripts, slides), improving student quiz scores by 35% in a structured pilot evaluation.
- Deployed Whisper ASR and ElevenLabs TTS into a WebRTC conversational UI, achieving sub-1.2s end-to-end latency for real-time voice Q&A at scale.

Project 990 Inc, Indiana | *AI Engineer* Jan 2026 – May 2026

- Led 2 data analysts to design a prompt-agnostic LLM classification pipeline to categorize 100K+ nonprofit mission statements into 27 NTEE codes using Llama 3.3 70B and Gemma by leveraging Tree-of-Thought reasoning.
- Accelerated identification of grant funding disparities across 3,000+ US counties from weeks to minutes for a 2.97M-record IRS 990 dataset by building a 9-module interactive analytics pipeline using Pandas, Plotly, and SciPy.

Digbi Health, California | *Data Science Intern* Jun 2025 – Nov 2025

- Shipped FastAPI REST endpoints backed by PostgreSQL to clean, normalize, and CPT-tag clinical datasets, reduced pre-processing runtime 40% using Python and SQL stratification logic.
- Productionized an MCP-based Text-to-SQL interface using LLM-driven query generation, enabling 15+ clinicians and staff to query claims data via natural language thereby eliminating SQL dependency entirely.

Edelweiss Global Markets, India | *Data Engineer* Jul 2022 – Aug 2024

- Orchestrated 40+ automated ETL pipelines using Airflow, PySpark, and Hadoop on AWS to process 1 TB/day of equities and futures data.
- Implemented a validation framework using Great Expectations and Polars on AWS S3 tick data, increased data integrity from 85% to 98% across 500K+ records/day.
- Streamlined tick-level 10-12 TB daily data cleanup post-validation via shell scripts on AWS S3, saved 8+ hours/week and improved storage utilization 30%.
- Engineered PySpark pipelines on Kubernetes to process 1M+ TBT market records daily, ingesting into ClickHouse for sub-second analytical queries; reduced TBT processing time by 12+ hours through optimized data partitioning and pod scaling.

ACADEMIC PROJECTS

Real-Time Trade Anomaly Detection Engine

- Architected a real-time trade anomaly detection pipeline using Kafka, PySpark Streaming, and an ensemble of Isolation Forest and XGBoost, achieving 91% precision and 88% recall on spoofing patterns.
- Optimized classification threshold via ROC and precision-recall analysis, reducing false positive rate by 18% while meeting regulatory-grade recall requirements.

MeetSmart – AI-Powered Meeting Companion

- Built a real-time meeting intelligence platform using FastAPI and Whisper ASR to process raw audio across 5+ languages via multilingual transcription and abstractive NLP summarization.
- Developed automated action-item and VADER sentiment extraction pipelines in Python, reducing post-meeting documentation time by 80% and improving action-tracking accuracy by 60%.

ACHIEVEMENTS

- 1st Place, Honeywell Hackathon @ NSBE 2026:** Built UnStutter AI, a real-time speech accessibility tool for 70M+ people worldwide affected by speech disfluencies; outperformed 35+ competing teams.
- 2nd Place, Claude Hackathon:** Architected CareBridge, a multi-agent AI platform automating post-discharge patient follow-ups to reduce hospital readmission rates.